

Practical Machine Learning - Prediction Assignment

Mehran Behzad

26 6 2020

Intro

Data from fitness tracking devices were gathered for subjects doing curls in 4 different manners (see <http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>). The goal of this assignment is to build a training model based on these data capable of predicting the type of exercise based on the movement data measured by these devices.

Data exploration and cleanup

```
# load libraries
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(caret)

## Loading required package: lattice
## Loading required package: ggplot2

# read data
training <- read.csv(file = 'data/pml-training.csv')
testing <- read.csv(file = 'data/pml-testing.csv')
```

The table has 19622 rows of data, each has 160 features.

```
#
head(training[(1:10)])

##   X user_name raw_timestamp_part_1 raw_timestamp_part_2  cvtd_timestamp
## 1 1 carlitos      1323084231          788290 05/12/2011 11:23
## 2 2 carlitos      1323084231          808298 05/12/2011 11:23
## 3 3 carlitos      1323084231          820366 05/12/2011 11:23
## 4 4 carlitos      1323084232          120339 05/12/2011 11:23
## 5 5 carlitos      1323084232          196328 05/12/2011 11:23
## 6 6 carlitos      1323084232          304277 05/12/2011 11:23
## new_window num_window roll_belt pitch_belt yaw_belt
## 1          no          11      1.41      8.07    -94.4
```

```
## 2      no      11      1.41      8.07     -94.4
## 3      no      11      1.42      8.07     -94.4
## 4      no      12      1.48      8.05     -94.4
## 5      no      12      1.48      8.07     -94.4
## 6      no      12      1.45      8.06     -94.4
```

It seems the first 7 columns are user specific or some time stamps and not task specific. We can remove them for the training.

```
# remove first columns which are not relevant for the exercise
training <- training[-(1:7)]
```

Furthermore, there are a lot of features that are almost N/A for all observations. We will remove them:

```
# remove column which have at least 90% NAs
training <- training[, apply(training, 2, function(y) length(which(is.na(y))) < .90 * dim(training)[1])
# same for test set
testing <- testing[-(1:7)]
testing <- testing[, apply(testing, 2, function(y) length(which(is.na(y))) < .90 * dim(testing)[1])]
```

We will also remove any feature/column which is missing in the Training or Testing data set which won't help with prediction or training.

```
# memorize the classe column which is deliberately missing in the testing dataset and has to be predicted
classe <- training$classe

# remove those which are only present in one dataset
testing <- testing[, intersect(colnames(training), colnames(testing))]
training <- training[, intersect(colnames(training), colnames(testing))]
training$classe <- classe
```

This gives us 53 features which we will feed the models with.

Training

Now we will split our training data, to allow testing its prediction with known observations:

```
inTrain <- createDataPartition(y= training$classe, p=0.7, list=FALSE)
trainingSub <- training[inTrain,]
testingSub <- training[-inTrain,]
```

Using the Random Forest (rf) algorithm to train the model

```
model_rf <- train(classe ~ ., method="rf", data = trainingSub,
                 trControl = trainControl(method = "cv"),
                 number = 3, na.action=na.exclude)
```

We build a second model using lda technique:

```
model_lda <- train(classe ~ ., method="lda", data = trainingSub,
                  trControl= trainControl(method="cv", number=10),
                  na.action=na.exclude)
```

Evaluation

```
confusionMatrix(testingSub$classe, predict(model_rf, testingSub))
```

```
## Confusion Matrix and Statistics
##
```

```
##           Reference
## Prediction    A    B    C    D    E
##           A 1674    0    0    0    0
##           B    6 1132    1    0    0
##           C    0    8 1018    0    0
##           D    0    0   13  950    1
##           E    0    0    1    2 1079
##
## Overall Statistics
##
##           Accuracy : 0.9946
##           95% CI : (0.9923, 0.9963)
##           No Information Rate : 0.2855
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9931
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9964  0.9930  0.9855  0.9979  0.9991
## Specificity      1.0000  0.9985  0.9984  0.9972  0.9994
## Pos Pred Value   1.0000  0.9939  0.9922  0.9855  0.9972
## Neg Pred Value   0.9986  0.9983  0.9969  0.9996  0.9998
## Prevalence       0.2855  0.1937  0.1755  0.1618  0.1835
## Detection Rate   0.2845  0.1924  0.1730  0.1614  0.1833
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9982  0.9958  0.9919  0.9975  0.9992
```

```
confusionMatrix(testingSub$classe, predict(model_lda, testingSub))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1389   36  120  126    3
##           B  182  743  133   37   44
##           C  119   99  672  116   20
##           D   69   42  108  705   40
##           E   52  192   93   95  650
##
## Overall Statistics
##
##           Accuracy : 0.7067
##           95% CI : (0.6949, 0.7183)
##           No Information Rate : 0.3077
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6282
##
## McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
```

```
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.7670   0.6682   0.5968   0.6534   0.8587
## Specificity      0.9300   0.9170   0.9256   0.9461   0.9158
## Pos Pred Value   0.8297   0.6523   0.6550   0.7313   0.6007
## Neg Pred Value   0.8998   0.9223   0.9066   0.9240   0.9777
## Prevalence       0.3077   0.1890   0.1913   0.1833   0.1286
## Detection Rate   0.2360   0.1263   0.1142   0.1198   0.1105
## Detection Prevalence 0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy 0.8485   0.7926   0.7612   0.7997   0.8872
```

It seems that the random forest method gives us a model with a high (~99%) accuracy.

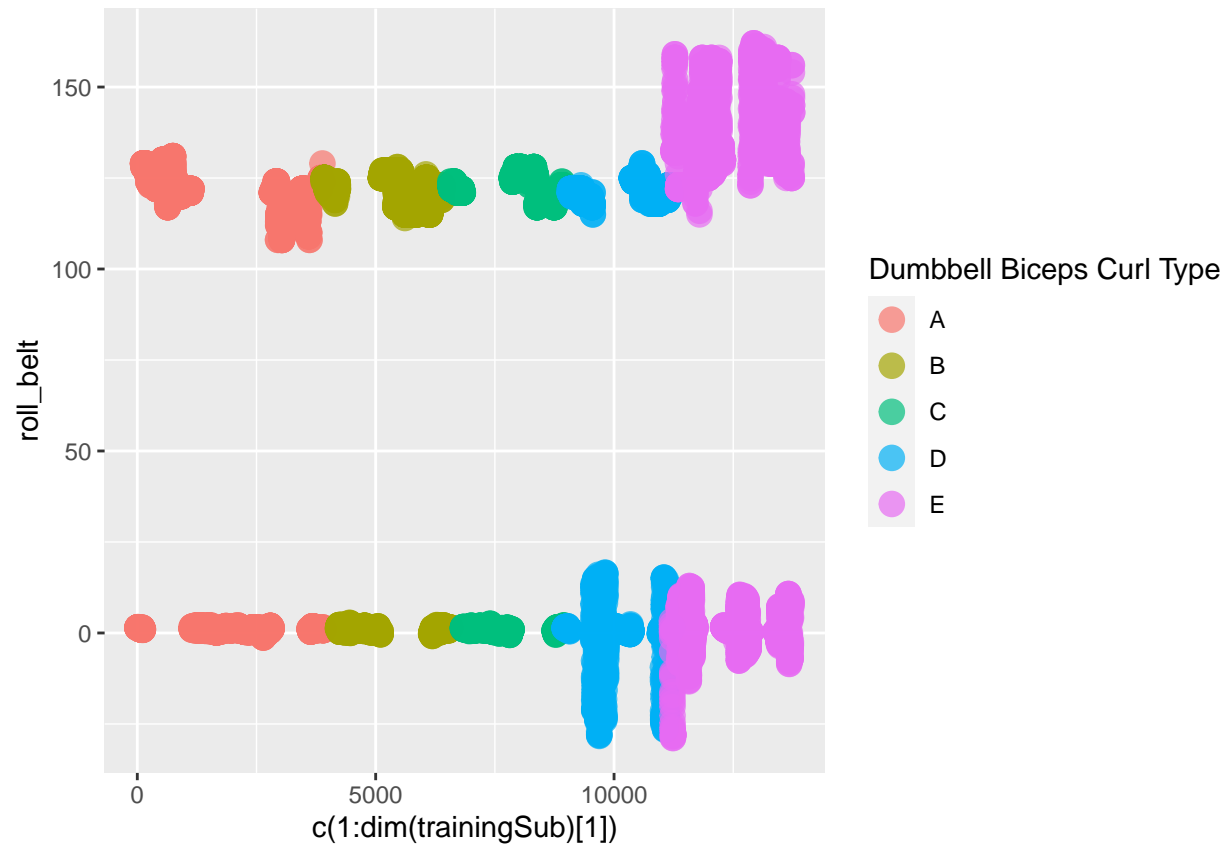
Let's see which parameters have the most influence:

```
varImp(model_rf)

## rf variable importance
##
##   only 20 most important variables shown (out of 52)
##
##               Overall
## roll_belt      100.00
## yaw_belt       79.70
## magnet_dumbbell_z 66.34
## magnet_dumbbell_y 60.41
## pitch_belt     59.82
## pitch_forearm  54.47
## magnet_dumbbell_x 52.32
## roll_forearm   50.07
## accel_belt_z   44.97
## magnet_belt_z  43.23
## accel_dumbbell_y 41.54
## roll_dumbbell  41.52
## magnet_belt_y  39.90
## accel_dumbbell_z 35.37
## roll_arm       33.25
## accel_forearm_x 32.90
## accel_arm_x    28.32
## total_accel_dumbbell 28.25
## accel_dumbbell_x 27.39
## yaw_dumbbell   27.26
```

And the one with most effect plotted:

```
ggplot(trainingSub, aes(y=roll_belt, x = c(1:dim(trainingSub)[1]), colour = factor(classe))) +
  geom_point(size=4, alpha=0.7) +
  labs(color = "Dumbbell Biceps Curl Type")
```



Predicting testing data

```
predict(model_rf, testing)
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```